

Operating Systems

I/O

Overview

- Devices and Controllers
- I/O Subsystem
- Device Drivers

I/O Controllers

- A device controller is attached to the system or integrated into the motherboard or SoC
- The peripheral itself attaches to the controller
 - RS-232, SCSI, SATA, SAS, USB
- Convert a string of bits into bytes or blocks of bytes
 - Even disks are strings of bits
- The controller has registers mapped into memory
 - Read and written to control and check status

I/O Ports vs Memory Mapping

- I/O ports in a dedicated namespace
- Accessed with special I/O instructions
 - `outb %al,$18`
- As opposed to memory space which is accessed with standard load, store, move

```
$ cat /proc/ioports
0000-0cf7 :PCI Bus
0000:00 0000-001f : dma1
0020-0021 : pic1
0040-0043 : timer0
0050-0053 : timer1
0060-0060 : keyboard
0064-0064 : keyboard
0070-0071 : rtc0
0080-008f : dma page reg
00a0-00a1 : pic2
00c0-00df : dma2
00f0-00ff : fpu
02f8-02ff : serial
0378-037a : parport0
03c0-03df : vga+
03f2-03f2 : floppy
```

The I/O Subsystem

- Devices have a complicated low-level interface
 - Control and data registers mapped into memory
 - Header files and documentation to understand registers and bit fields
- Goal: provide a high level interface so that programs don't have to be rewritten
 - I/O devices do mostly the same thing – they input, they output
 - Design: a small set of of abstract routines can encapsulate various devices

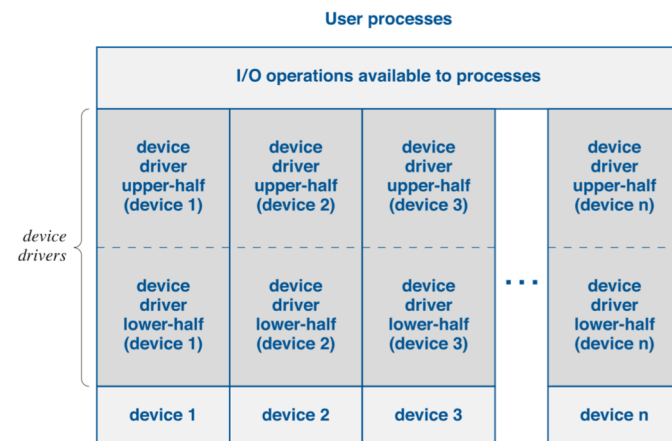
I/O Interfaces

- Another purpose of the I/O interface is to protect shared I/O resources (devices, buffers)
- Safe and fair access
- Policies can be applied at the high level interfaces and can be generalized over various devices
- The Unix abstraction is that “everything is a file”

Organization of the I/O Subsystem

- High level I/O functions to abstract the details of hardware and provide general entry points
 - Design challenge is to capture diversity of devices with a generic interface
- Device drivers interact with specific devices
- Drivers have a upper and lower half
- The upper half interacts with process requests
- The lower half responds to interrupts with handler functions
 - service interrupts, initiate new operations as necessary

Organization of the I/O Subsystem



I/O and Driver Abstractions

- Synchronous vs Asynchronous
 - Synchronous: the requesting process is blocked until I/O completes – easier to program
 - Asynchronous: the process can continue to execute – more control of overlap of communication and computation
- Asynchronous I/O interfaces must notify the process
 - Deliver a signal
 - Spawn a thread
 - Check (poll) the status of an I/O action, or read from a queue of completion actions

POSIX AIO

- Allows applications to initiate one or more asynchronous I/O operations
 - signal, instantiate a thread, no notification
- `aio_read(struct aiocb *aiocbp)`, `aio_write()`
- `aio_return()` – to check the return status of an AIO operation

```
struct aiocb { /* The order of these fields is implementation-dependent */
    int aio_fildes; /* File descriptor */
    off_t aio_offset; /* File offset */
    volatile void *aio_buf; /* Location of buffer */
    size_t aio_nbytes; /* Length of transfer */
    int aio_reqprio; /* Request priority */
    struct sigevent aio_sigevent; /* Notification method */
    .....
};
```

I/O and Driver Abstractions

- Format and size of data transfers
 - Bytes, strings, blocks
- Block vs character interfaces
 - Look in `/dev` on your favorite *nix system
 - The question is whether chunks of data (blocks) are independently addressable
- Buffering can be used to adapt between the two
- How much state is preserved between requests?
 - Specify a starting point and read successive blocks or specify a block with each read?

Abstract I/O Interface

Operation	Purpose
<code>close</code>	Terminate use of a device
<code>control</code>	Perform operations other than data transfer
<code>getc</code>	Input a single byte of data
<code>init</code>	Initialize the device at system startup
<code>open</code>	Prepare the device for use
<code>putc</code>	Output a single byte of data
<code>read</code>	Input multiple bytes of data
<code>seek</code>	Move to specific data (usually a disk)
<code>write</code>	Output multiple bytes of data

Open, Read, Write, Close

- Common paradigm (Xinu, Unix, Windows)
- Before a process can use a device, it must open it
 - Manage exclusive access
 - Check permissions
 - Set up state in system data structures
- Close when finished
 - Clean up state
 - The device could be powered down

Control

- Control interface allows for configuration of device driver parameters
- Can also manage device-specific interactions that are not possible with the standard interfaces
 - Buffering or caching behavior
- `ioctl()` on Unix

Binding Operations and Devices

- Abstract interfaces need to act on specific devices
- Must be mapped to device driver functions
- The OS provides a virtual I/O environment, passing operations through to devices via drivers
- Unix embeds devices in the filesystem, providing names to specific devices
- General-purpose OSes construct this dynamically, but embedded systems often statically configure it

Device Names in Xinu

- Specify a set of devices when the system is configured
- Assign an integer device descriptor
- For instance, `CONSOLE` is device 0
- Programs don't need to be rewritten when devices change, but the system does need to be reconfigured and recompiled

Xinu's Device Switch Table

- The OS must forward I/O operations to the correct driver function
- The device ID is used as an index into a table of device-specific functions
- Each entry in the table contains information about the device and function pointers to functions that implement operations
- To write to a device, find the device entry and invoke the specific write function
- Xinu's approach is simple but is fundamentally the same as e.g. Unix

Xinu Example

	open	close	read	write	getc	
CONSOLE	conopen	conclose	conread	conwrite	congetc	
ETHER	ethopen	ethclose	ethread	ethwrite	ethgetc	...
DISK	dskopen	dskclose	dskread	dskwrite	dskgetc	
			⋮			

- Uniform interface hiding the differences of underlying hardware

Multiple Instances of a Device

- Multiple instances of a device can share a driver
- Multiple instances in the device table that are largely the same, differing in only a few aspects
- Each instance will have its own control and status registers
- Can also be distinguished by the “minor” device number

Device Table Entry

```
/* Device table entry */
struct dentry {
    int32  dvnum;
    int32  dvminor;
    char   *dvname;
    devcall (*dvinit) (struct dentry *);
    devcall (*dvopen) (struct dentry *, char *, char *);
    devcall (*dvclose) (struct dentry *);
    devcall (*dvread) (struct dentry *, void *, uint32);
    devcall (*dvwrite) (struct dentry *, void *, uint32);
    devcall (*dvseek) (struct dentry *, int32);
    devcall (*dvgetc) (struct dentry *);
    devcall (*dvputc) (struct dentry *, char);
    devcall (*dvctl) (struct dentry *, int32, int32, int32);
    void    *dvcsr;
    void    (*dvintr) (void);
    byte    dvirq;
};
```

Some Devices

```
extern struct dentry devtab[]; /* one entry per device */

/* Device name definitions */

#define CONSOLE    0      /* type tty      */
#define NULLDEV    1      /* type null     */
#define ETHER0     2      /* type eth      */
#define NAMESPACE 3      /* type nam      */
#define RDISK      4      /* type rds      */
#define RAM0       5      /* type ram      */
#define RFILESYS   6      /* type rfs      */
#define RFILE0     7      /* type rfl      */
#define RFILE1     8      /* type rfl      */
#define RFILE2     9      /* type rfl      */
#define RFILE3    10      /* type rfl      */
#define RFILE4    11      /* type rfl      */
#define RFILE5    12      /* type rfl      */
#define RFILE6    13      /* type rfl      */
```

read()

```
syscall read(
    did32 descrp, /* Descriptor for device */
    char *buffer, /* Address of buffer */
    uint32 count /* Length of buffer */
)
{
    intmask mask; /* Saved interrupt mask */
    struct dentry *devptr; /* Entry in device switch table */
    int32 retval; /* Value to return to caller */

    mask = disable();
    if (isbaddev(descrp)) {
        restore(mask);
        return SYSERR;
    }
    devptr = (struct dentry *) &devtab[descrp];
    retval = (*devptr->dvread) (devptr, buffer, count);
    restore(mask);
    return retval;
}
```

control()

```
syscall control(
    did32 descrp, /* Descriptor for device */
    int32 func, /* Specific control function */
    int32 arg1, /* Specific argument for func */
    int32 arg2 /* Specific argument for func */
)
{
    intmask mask; /* Saved interrupt mask */
    struct dentry *devptr; /* Entry in device switch table */
    int32 retval; /* Value to return to caller */

    mask = disable();
    if (isbaddev(descrp)) {
        restore(mask);
        return SYSERR;
    }
    devptr = (struct dentry *) &devtab[descrp];
    retval = (*devptr->dvcntl) (devptr, func, arg1, arg2);
    restore(mask);
    return retval;
}
```

open() and close()

- Implemented identically to read() and control()
- Explicit opening and closing allows the system to maintain a *reference count* of processes using a device
 - Again, the system may power a device down when not in use

Null Entries in Devtab

- Note that each of the high-level functions calls the device-specific function without checking its validity
- Not all operations make sense on all devices
 - You can't seek on the console, or getc() on a network device
- ionull() returns OK
- ioerr() returns SYSERR

Initialization

- General operating systems can dynamically initialize devices
 - recall the discussion of USB devices
- Embedded systems like Xinu use static configuration
- Xinu specifies devices and functions in a file called Configuration, and generates a C file and a header with appropriate values

Configuration

```
/* Entries for a device specify the functions that handle each of the */
/* high-level I/O operations as follows: */
/*
    */
/* -i  init  -o  open  -c  close  */
/* -r  read  -w  write -s  seek   */
/* -g  getc  -p  putc  -n  control */
/* -intr int_hndlr -csr csr -irq irq */
/*
    */
/*****

/* Type Declarations for both real- and pseudo- device types */

/* type of a null device */
null:
    on nothing
    -i ionull -o ionull -c ionull
    -r ionull -g ionull -p ionull
    -w ionull -s ioerr
```

Configuration

```
/* type of a tty device */
tty:
    on uart
        -i ttyinit      -o ionull      -c ionull
        -r ttyread      -g ttygetc     -p ttyputc
        -w ttywrite     -s ioerr       -n ttycontrol
        -intr ttyhandler

/* type of a ethernet device */
eth:
    on am335x_eth
        -i ethinit      -o ioerr       -c ioerr
        -r ethread      -g ioerr       -p ioerr
        -w ethwrite     -s ioerr       -n ethcontrol
        -intr ethhandler
```

conf.c

```
struct dentry devtab[NDEVS] =
{
/**
 * Format of entries is:
 * dev-number, minor-number, dev-name,
 * init, open, close,
 * read, write, seek,
 * getc, putc, control,
 * dev-csr-address, intr-handler, irq
 */

/* CONSOLE is tty */
{ 0, 0, "CONSOLE",
  (void *)ttyinit, (void *)ionull, (void *)ionull,
  (void *)ttyread, (void *)ttywrite, (void *)ioerr,
  (void *)ttygetc, (void *)ttyputc, (void *)ttycontrol,
  (void *)0x44e09000, (void *)ttyhandler, 72 },
```

conf.c

```
/* CONSOLE is tty */
{ 0, 0, "CONSOLE",
  (void *)ttyinit, (void *)ionull, (void *)ionull,
  (void *)ttyread, (void *)ttywrite, (void *)ioerr,
  (void *)ttygetc, (void *)ttyputc, (void *)ttycontrol,
  (void *)0x44e09000, (void *)ttyhandler, 72 },

/* NULLDEV is null */
{ 1, 0, "NULLDEV",
  (void *)ionull, (void *)ionull, (void *)ionull,
  (void *)ionull, (void *)ionull, (void *)ioerr,
  (void *)ionull, (void *)ionull, (void *)ioerr,
  (void *)0x0, (void *)ioerr, 0 },

/* ETHER0 is eth */
{ 2, 0, "ETHER0",
  (void *)ethinit, (void *)ioerr, (void *)ioerr,
  (void *)ethread, (void *)ethwrite, (void *)ioerr,
  (void *)ioerr, (void *)ioerr, (void *)ethcontrol,
  (void *)0x0, (void *)ethhandler, 0 },
```

Embedded Linux - DeviceTree

- Embedded Linux uses something called DeviceTree
- A Flattened Device Tree (FDT) is shipped so that a kernel image can be configured appropriately for hardware
- Required for Linux on new ARM SoCs
- Identifies the type of CPU and describes devices in the system very similarly to what we have discussed

Summary

- Complex operating systems have more functionality between the abstract interfaces and the device drivers
- Caching
- Security and policy
- This indirection mechanism forms the basis however